

COSC3500 Tutorial 3: Experiments in matrix multiplication optimisation

17th August 2021

Matrix Multiply(1): naive method

```
double A[n][n];
double B[n][n];
double C[n][n];

int trials = atoi(argv[1]);
for (int i = 0; i < n; i++) {
    for (int j = 0; j < n; j++) {
        A[i][j] = (double) rand() / (double) RAND_MAX;
        B[i][j] = (double) rand() / (double) RAND_MAX;
        C[i][j] = 0;
    }
}

std::vector<int> times;
for (int i = 0; i < trials; i++) {
    auto start = std::chrono::high_resolution_clock::now();
    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            for (int k = 0; k < n; k++) {
                C[i][j] += A[i][k] * B[k][j];
            }
        }
    }
    auto stop = std::chrono::high_resolution_clock::now();
    auto duration = std::chrono::duration_cast<std::chrono::microseconds>(stop - start).count();
    times.push_back(duration);
}
```

*Trials done
at n = 2048*

Runtime

Version	Method	Timing
1	Standard Implementation	242.06s

Matrix Multiply(2): Switch inner loops

```
for (int i = 0; i < n; i++) {  
    for (int k = 0; k < n; k++) {  
        for (int j = 0; j < n; j++) {  
            C[i][j] += A[i][k] * B[k][j];  
        }  
    }  
}
```

Runtime


Version	Method	Timing
1	Standard Implementation	242.06s
2	Interchange Loop Order	93.20s

Remember, C uses row-major order for multi-dimensional arrays: first index is row, second index is column!

E.g. `arr[i][j]` indexes row `i`, column `j`.

Spatial Locality case 1


$$C[i][j] += A[i][k] * B[k][j]$$
$$i = 1, j = 2, k = 0..n$$



Spatial Locality case 2

$$C[i][j] += A[i][k] * B[k][j]$$

$i = 1, k = 2, j = 0..n$



Valgrind

Usage: *valgrind --tool=cachegrind args*

Example: `valgrind --tool=cachegrind ./matrix 1`

```
==251==
==251== D   refs:      66,180,043,483 (61,391,524,032 rd + 4,788,519,451 wr)
==251== D1  misses:    4,718,855,645 ( 4,717,280,497 rd +      1,575,148 wr)
==251== LLd misses:    4,710,839,645 ( 4,709,265,294 rd +      1,574,351 wr)
==251== D1  miss rate:      7.1% (      7.7% +      0.0% )
==251== LLd miss rate:      7.1% (      7.7% +      0.0% )
==251==
==251== LL refs:      4,718,857,147 ( 4,717,281,999 rd +      1,575,148 wr)
==251== LL misses:    4,710,841,133 ( 4,709,266,782 rd +      1,574,351 wr)
==251== LL miss rate:      1.8% (      1.8% +      0.0% )
```

Modified Loop Order

```
==259== D   refs:      28,503,779,515 (26,407,360,706 rd + 2,096,418,809 wr)
==259== D1  misses:    253,851,805 ( 252,276,657 rd + 1,575,148 wr)
==259== LLd misses:    253,776,053 ( 252,201,702 rd + 1,574,351 wr)
==259== D1  miss rate: 0.9% ( 1.0% + 0.1% )
==259== LLd miss rate: 0.9% ( 1.0% + 0.1% )
==259==
==259== LL refs:      253,853,307 ( 252,278,159 rd + 1,575,148 wr)
==259== LL misses:    253,777,541 ( 252,203,190 rd + 1,574,351 wr)
==259== LL miss rate: 0.2% ( 0.2% + 0.1% )
```

Runtime

Version	Method	Timing
1	Standard Implementation	242.06s
2	Interchange Loop Order	93.20s
3	O3 Optimisation Flag	4.31s

Why is `-O3` so much faster?

- `-O3` flag enables vectorisation by default
- Vectorisation is very powerful and the compiler could easily vectorise this code. Sometimes vectorisation is not possible or needs to be done manually

Matrix Multiply (3): blocking

```
for (int ii = 0; ii < n; ii += b) {  
    for (int jj = 0; jj < n; jj += b) {  
        for (int kk = 0; kk < n; kk += b) {  
            for (int i = 0; i < b; i++) {  
                for (int k = 0; k < b; k++) {  
                    for (int j = 0; j < b; j++) {  
                        C[i + ii][j + jj] += A[i + ii][k + kk] * B[k + kk][j + jj];  
                    }  
                }  
            }  
        }  
    }  
}
```

Animations of the different approaches:

Method 1 - Naive: <https://www.youtube.com/watch?v=QYpH-847z0E>

Method 2 - Loop interchange: https://www.youtube.com/watch?v=0u2K_dRLhWw

Method 3 - Blocking with loop interchange: <https://www.youtube.com/watch?v=aMvCEEBIBto>

Runtime

Version	Method	Timing
1	Standard Implementation	242.06s
2	Interchange Loop Order	93.20s
3	O3 Optimisation Flag	4.31s
4	More Cache Optimisation	2.73s

Summary of Methods

- Utilising computer hardware well (Show levels of cache)
- Compiler Flags
- Achieved ~5 times speedup with designing around caching
- Experiment a lot! (We can still do better)